

IBS · AGENTIC AI PRACTITIONER PROGRAM

Day 1.5 · Mid-session exercise

Pick the Model

Five scenarios. No right answer. One real question.

WHY YOU'RE DOING THIS

You have just spent 90 minutes on tokens, parameters, families, variants, and the open-vs-closed trade-off. The temptation now is to memorize the slides. Don't. The only thing that matters in your actual job is whether you can **look at a real-world scenario and pick something defensible**. These five scenarios force exactly that.

HOW IT WORKS

- **Your pair gets ONE scenario.**
Don't read the others first — work yours cold.
- **Spend 6 minutes on it.**
Don't aim for a perfect answer. Aim for a defensible one.
- **Fill every field on your page.**
Especially the last one. That's where the real learning happens.
- **Be ready to defend your pick.**
Group debrief: 5 minutes. Three patterns will surface.

BEFORE YOU PICK A MODEL — THE FIVE QUESTIONS YOU MUST ASK

1. What's the volume? (per day / week / month)
2. What languages? (English-only? Multilingual? Code-switched?)
3. What's the data residency constraint? (Can data leave AWS? The country? The customer's tenant?)
4. Is this latency-sensitive (sub-second) or batch-tolerant (hours OK)?
5. What's the consequence of a wrong answer? (Regulatory? Financial? Annoying?)

If you skip these five questions and jump straight to picking a model — that's exactly the failure mode this exercise is built to expose.



SCENARIO

The 24×7 chatbot

“

*IBS is building a customer-facing chatbot for an airline. It answers around **200,000 questions per day** across English, Hindi, Tamil, and Mandarin. Most are FAQs — "what's my baggage allowance," "can I cancel for free?" About 5% are complex enough that a wrong answer would land the airline in front of a consumer-rights regulator. Latency target: under 2 seconds end-to-end. **Pick a model.***

STEP 1 · ANSWER THESE BEFORE YOU PICK ANYTHING

Q1 VOLUME

per day / week / month

Q2 LANGUAGES

English only? mixed?

Q3 DATA RESIDENCY

what's the constraint?

Q4 LATENCY

real-time? batch?

Q5 WRONG-ANSWER COST

regulatory? financial?

STEP 2 · YOUR RECOMMENDATION

MODEL PICK + TIER

family · model · tier · and why

MONTHLY COST

rough \$ / month

STEP 3 · THE FIELD THAT MATTERS MOST

Name the **ONE** assumption that, if wrong, would completely change your answer.

THE ONE ASSUMPTION

B

SCENARIO

The contract reviewer

“

An IBS travel-insurance underwriter wants to summarize **80-page policy contracts** and flag clauses that deviate from the standard template. Volume is small — **50 contracts per week**. The CFO has said in writing that no contract content may leave AWS infrastructure, ever. The underwriter is fine waiting an hour for a result. **Pick a model.**

STEP 1 · ANSWER THESE BEFORE YOU PICK ANYTHING

Q1 VOLUME

per day / week / month

Q2 LANGUAGES

English only? mixed?

Q3 DATA RESIDENCY

what's the constraint?

Q4 LATENCY

real-time? batch?

Q5 WRONG-ANSWER COST

regulatory? financial?

STEP 2 · YOUR RECOMMENDATION

MODEL PICK + TIER

family · model · tier · and why

MONTHLY COST

rough \$ / month

STEP 3 · THE FIELD THAT MATTERS MOST

Name the **ONE** assumption that, if wrong, would completely change your answer.

THE ONE ASSUMPTION

C

SCENARIO

The intern's experiment

“

*A junior dev on your team wants to spend a weekend building a prototype: a tool that auto-generates Python unit tests from a function signature. They plan to run it about **500 times** across the team's codebase to see how the output looks. There's no production traffic. There's no budget approval. There's no SLA. **Pick a model.***

STEP 1 · ANSWER THESE BEFORE YOU PICK ANYTHING

Q1 VOLUME

per day / week / month

Q2 LANGUAGES

English only? mixed?

Q3 DATA RESIDENCY

what's the constraint?

Q4 LATENCY

real-time? batch?

Q5 WRONG-ANSWER COST

regulatory? financial?

STEP 2 · YOUR RECOMMENDATION

MODEL PICK + TIER

family · model · tier · and why

MONTHLY COST

rough \$ / month

STEP 3 · THE FIELD THAT MATTERS MOST

Name the **ONE** assumption that, if wrong, would completely change your answer.

THE ONE ASSUMPTION

D

SCENARIO

The Mumbai transcriber

“

*IBS runs an outsourced call centre in Mumbai handling **8,000 hours of recorded calls per day**. The ops team wants to transcribe each call, summarize it, and extract action items into a ticketing system. The calls are in **Hinglish** — **agents and customers switch between Hindi and English mid-sentence**. Calls don't need to be processed in real time; overnight batch is fine. **Pick a model.***

STEP 1 · ANSWER THESE BEFORE YOU PICK ANYTHING

Q1 VOLUME

per day / week / month

Q2 LANGUAGES

English only? mixed?

Q3 DATA RESIDENCY

what's the constraint?

Q4 LATENCY

real-time? batch?

Q5 WRONG-ANSWER COST

regulatory? financial?

STEP 2 · YOUR RECOMMENDATION

MODEL PICK + TIER

family · model · tier · and why

MONTHLY COST

rough \$ / month

STEP 3 · THE FIELD THAT MATTERS MOST

Name the **ONE** assumption that, if wrong, would completely change your answer.

THE ONE ASSUMPTION

E

SCENARIO

The standardization mandate

“

*IBS leadership has asked you to recommend **one single model family** that the entire company will standardize on for the next 18 months. They want "one throat to choke," consistent pricing, and a single set of prompt patterns. They've asked for your recommendation by end of week. **Pick a model.***

STEP 1 · ANSWER THESE BEFORE YOU PICK ANYTHING

Q1 VOLUME

per day / week / month

Q2 LANGUAGES

English only? mixed?

Q3 DATA RESIDENCY

what's the constraint?

Q4 LATENCY

real-time? batch?

Q5 WRONG-ANSWER COST

regulatory? financial?

STEP 2 · YOUR RECOMMENDATION

MODEL PICK + TIER

family · model · tier · and why

MONTHLY COST

rough \$ / month

STEP 3 · THE FIELD THAT MATTERS MOST

Name the **ONE** assumption that, if wrong, would completely change your answer.

THE ONE ASSUMPTION

AFTER THE GROUP DEBRIEF

What you actually learned

01

Which of the five reflex questions did you skip?

*Most teams skip the data residency question or the volume question. Both flip the answer.***YOUR NOTES**

02

Were you about to pick a tier (Opus / Sonnet / Haiku) based on task difficulty?

*Tier should follow volume, not difficulty. Cheap models on hard tasks at low volume is often the right call.***YOUR NOTES**

03

What did you assume that you can't actually verify until you ask the client?

*Capture this. It's the start of your discovery question list for any real Bedrock engagement.***YOUR NOTES**

WHERE THIS LEADS NEXT

Several of your scenarios — especially the one where data couldn't leave AWS, and the one where you wanted to standardize — point to a platform-shaped answer, not a model-shaped one. That's the bridge into the next section. **If you can call any model and the platform handles residency, IAM, audit, and routing — you've stopped picking models and started picking architectures.**