

Build a PRD. Then build a harder one.

Two scenarios. The first has obvious answers. The second has trade-offs that fight each other. The skill is not answering all eight questions — it is naming the conflicts and resolving them on paper.

THE EIGHT QUESTIONS

#	QUESTION	WHAT IT FORCES YOU TO DECIDE
1	Who is the user?	Persona, surface, channel. Internal vs external matters.
2	What is the goal?	One sentence. If it takes three, it has scope creep.
3	Success KPI	Measurable. Number + time window. "Better experience" is not a KPI.
4	Autonomy boundary	What's free, what needs approval, what's prohibited. Name operations.
5	What does failure look like?	Concrete worst-cases. Define the unacceptable, then design against it.
6	Rollback path	What reverses on partial failure? Manual or automated?
7	PII surface	Where PII enters, how it's masked, retention policy.
8	Cost ceiling	Per session, per day, per month. Above this, halt or escalate.

HOW THIS EXERCISE WORKS

Part 1 (20 min): build the easy PRD. You will get one defensible answer per question with light effort. ChatGPT helps; you sanity-check.

Part 2 (25–35 min): build the tricky PRD. Three of the eight questions have answers that conflict with other questions. You must name each trade-off, pick a side with stated reasoning, and document what stays unresolved.

Deliverable: one ChatGPT conversation with both PRDs and your three trade-off resolutions. The tricky PRD is what your peers will review — not the easy one.

Part 1 — The easy PRD

Internal-only. Single user type. Read-only. Low blast radius. Every question has one defensible answer.

THE SCENARIO

"Build us a Slack assistant that summarises long threads on demand. A team member invokes it with /summarise on any thread, and it returns a 5-bullet summary plus a list of open questions, posted back as a reply only the invoker can see."

1 Run the easy PRD (15 min)

Open ChatGPT. Paste this prompt. Watch how cleanly it answers — there are no real conflicts in this brief.

PROMPT TO PASTE

I am drafting a PRD for an AI system. Here is the brief:

"A Slack assistant. A user invokes /summarise on any thread. It returns a 5-bullet summary plus a list of open questions, posted back as a reply only the invoker can see. Internal use only."

Answer all eight PRD questions in a numbered list. One sentence per answer. Be concrete.

1. Who is the user?
2. What is the goal? (one sentence)
3. Success KPI (number + time window, not "better experience")
4. Autonomy boundary (free / approval-needed / prohibited)
5. What does failure look like? (concrete)
6. Rollback path on partial failure
7. PII surface (where it enters, how it's handled)
8. Cost ceiling (per session, per day)

Rules:

- One sentence per answer.
- The KPI must be measurable.
- The autonomy boundary must name actual operations, not categories.

WHAT GOOD LOOKS LIKE

Every answer is one sentence and uncontroversial. The autonomy boundary is wide (read-only is safe). The PII surface is small (only Slack content the invoker can already see). The cost ceiling is generous. **If you find yourself debating answers, you are overthinking it — the easy PRD is supposed to feel easy.**

2 Verify the easy PRD (3 min)

Tick each box before moving on.

- ☐ All eight answers are exactly one sentence.
- ☐ The KPI has a number AND a time window (e.g., "80% of summaries marked useful within 7 days").
- ☐ The autonomy boundary names the actual operation ("read thread, post reply"), not a category ("helpful actions").
- ☐ Failure mode is concrete ("summary omits a critical decision"), not generic ("wrong output").
- ☐ PII surface is identified by data type, not just "some PII may be involved."
- ☐ Cost ceiling has a number (e.g., "\$0.05 per summary, halt above \$50/day").

Part 2 — The tricky PRD

Customer-facing. Multiple stakeholders. Real money moves. The eight questions stop being independent — answering one constrains another.

THE SCENARIO

"Build a customer-facing refund agent on our travel app. When a passenger reports a disruption (cancelled flight, missed connection, downgrade), the agent verifies the claim against airline data, computes the refund or credit per our policy, and processes it. Support cost per ticket is killing margins; CX wants instant resolution; finance wants every refund auditable; compliance flagged that PII handling will be reviewed."

Stakeholders and their stated wants:

Support ops: resolve 70% of disruption tickets without a human agent. Reduce average handle time from 14 minutes to under 60 seconds.

Finance: no refund above \$200 issued without a human approving it. Every refund needs an audit trail showing the policy clause invoked.

Customer experience: the customer should see a resolution in the chat session itself — no "we'll get back to you in 48 hours."

Compliance: passport numbers and payment method details appear in some tickets. Cannot be logged. Cannot leave the EU.

3 Read the four trade-offs before drafting (5 min)

Before you ask ChatGPT anything, read all four trade-offs on the next page. **Two of the eight PRD questions are battlegrounds, not questions.** If you draft the PRD without naming these first, you will write answers that contradict each other and not notice.

TRADE-OFF #1

Autonomy boundary vs Success KPI

STANCE A — Finance

No refund > \$200 without human approval. This is the only way to keep the audit clean and bound the blast radius of a misclassification.

STANCE B — Support ops

70% auto-resolution requires the agent to act alone on most claims. If every refund above \$200 needs human approval, auto-resolution rate collapses because the average disruption refund is \$340.

YOU MUST DECIDE

Where does the approval gate actually sit? Pick a threshold AND justify why it lets you hit the KPI without breaching the audit constraint. Name the policy clauses that the agent can act on alone.

TRADE-OFF #2

PII surface vs Customer experience

STANCE A — Compliance

Passport and payment data cannot be logged, cannot leave the EU, and cannot appear in retrieved context. The model must never see raw PII.

STANCE B — CX

Customers paste passport numbers and card details into the chat all the time. If the agent says "I cannot read that, please remove it," the experience is broken and ops costs spike from re-engagement.

YOU MUST DECIDE

How does the agent handle inbound messages that contain PII it must not see? Name the masking, the user-visible behaviour, AND who handles the tickets that cannot be processed safely.

TRADE-OFF #3

Failure mode vs Cost ceiling

STANCE A — Finance

Worst case is issuing a refund the customer was not entitled to. Mitigation: rigorous verification — check airline data, check policy, check duplicate claims. This burns tokens.

STANCE B — Engineering

Verification chain across 3 systems plus eval-as-judge plus retries on flaky upstream pushes per-session cost into the \$4–6 range. At expected volume that exceeds the entire support budget.

YOU MUST DECIDE

Which verifications run on every claim, which run only above a threshold, and what is the per-session cost ceiling? Defend the cost-vs-failure trade-off — there is no way to have both.

TRADE-OFF #4

Rollback path vs Customer experience

STANCE A — Finance

If a refund is mis-issued, it must be reversible. That means a 2-step process: agent proposes refund, refund queues for 5 minutes, then settles. The 5-minute window allows rollback.

STANCE B — CX

Customers see "refund processed" in 60 seconds. A queue introduces latency the experience metric won't tolerate. Anything other than instant feels like the old support flow.

YOU MUST DECIDE

Is there a third option? (Hint: instant for low-risk, queued for high-risk.) If you split the path, what is the threshold and what does the customer see in each branch?

4 Draft the tricky PRD WITH the trade-offs resolved (15 min)

Now draft the PRD. The prompt forces ChatGPT to address each trade-off explicitly rather than glossing over them.

PROMPT TO PASTE

I am drafting a PRD for a customer-facing refund agent on a travel app. The system has four known trade-offs that I have identified. I will list them, then ask you to write the PRD in a way that explicitly resolves each one.

Stakeholders:

- Support ops: 70% auto-resolution, <60s handle time
- Finance: refunds > \$200 need human approval, full audit trail
- CX: instant in-chat resolution
- Compliance: no PII logged, no PII out of EU, model never sees raw PII

Trade-offs to resolve:

1. Autonomy boundary vs auto-resolution KPI
2. PII surface vs CX (customers paste PII)
3. Failure-mode rigour vs per-session cost
4. Rollback path vs in-chat speed

Write the PRD answering all 8 questions. Rules:

- For Q3 (KPI), Q4 (autonomy), Q5 (failure), Q6 (rollback), Q7 (PII), Q8 (cost), explicitly state which trade-off you are resolving and how.
- For each trade-off, name the SIDE you are picking AND what you are giving up.
- Where a true compromise exists (e.g., split the path by risk), use it. Where one side must win, say so and justify with a number.
- At the bottom, add a "What stays unresolved" section listing anything you could not decide on paper — these go to a human stakeholder meeting on Monday.

HOW TO RESOLVE A TRADE-OFF — THE FOUR-STEP CHECK

- 1. NAME both sides.** Both stakeholders have a real reason. If you cannot articulate the opposing case in one sentence, you have not understood it.
- 2. QUANTIFY the cost of each side.** What does picking side A cost in concrete terms? Auto-resolution drops from 70% to 35%? Per-session cost rises from \$0.40 to \$2.10? Latency goes from 60s to 4 minutes? Numbers, not adjectives.
- 3. SEEK a structured compromise.** The best resolutions are rarely "pick one side." They are usually "split by a threshold" — low-risk takes the fast path, high-risk takes the safe path. Find the threshold that makes both sides bearable.
- 4. NAME what you gave up.** Every resolution sacrifices something. Write down what — explicitly. The PRD reader needs to see the cost of the decision, not just the decision.

5 Verify the tricky PRD (5 min)

This checklist is different from Part 1. It checks whether you made the trade-offs explicit, not just whether you answered.

- ☐ Each of the four trade-offs is named in the PRD by number (1, 2, 3, 4) — not implied.
- ☐ For each trade-off, the PRD states which side won AND what was given up.
- ☐ At least one trade-off uses a threshold-split resolution (not just "pick one side").
- ☐ The KPI (Q3) and the autonomy boundary (Q4) explicitly reference each other — they are linked, not independent.
- ☐ The PII handling (Q7) names the user-visible behaviour, not just the technical control.
- ☐ The cost ceiling (Q8) is justified against the failure-mode rigour from (Q5).
- ☐ The PRD ends with a "What stays unresolved" section. If yours has zero items there, you ducked a trade-off.
- ☐ The autonomy boundary names specific dollar thresholds AND specific policy clauses, not categories.

WHAT GOOD LOOKS LIKE FOR PART 2

A real practitioner-grade tricky PRD has 2–4 explicit trade-off resolutions in the body AND a non-empty "unresolved" list at the bottom. The unresolved list is not a failure — it is the artefact that goes to your VP / steering committee. The PRD that claims to resolve everything on paper is the PRD that fails in implementation.

HOW TO PRESENT YOUR TRICKY PRD

30-second pitch format: "For the refund agent, the four trade-offs are X, Y, Z, W. I resolved them by [picking side / splitting by threshold / escalating]. I gave up [thing]. What I could not resolve on paper: [list]." If your peers can poke a hole in the resolution, the PRD is doing its job — that is exactly what review is for.